

CM146, Winter 2018
Problem Set 4: Boosting, Multi-class Classification
Due March 8, 2018, 11:59pm

Submission instructions

- Submit your solutions electronically on the course Gradescope site as PDF files.
- If you plan to typeset your solutions, please use the LaTeX solution template. If you must submit scanned handwritten solutions, please use a black pen on blank white paper and a high-quality scanner app.

1 Boosting - 40 points

Consider the following examples $(x, y) \in \mathbb{R}^2$ (i is the example index):

i	x	y	Label
1	0	8	−
2	1	4	−
3	3	7	+
4	-2	1	−
5	-1	13	−
6	9	11	−
7	12	7	+
8	-7	-1	−
9	-3	12	+
10	5	9	+

In this problem, you will use Boosting to learn a hidden Boolean function from this set of examples. We will use two rounds of AdaBoost to learn a hypothesis for this data set. In each round, AdaBoost chooses a weak learner that minimizes the error ϵ . As weak learners, use hypotheses of the form (a) $f_1 \equiv [x > \theta_x]$ or (b) $f_2 \equiv [y > \theta_y]$, for some integers θ_x, θ_y (either one of the two forms, not a disjunction of the two). There should be no need to try many values of θ_x, θ_y ; appropriate values should be clear from the data. When using log, use base 2.

- (a) **[10 points]** Start the first round with a uniform distribution D_0 . Place the value for D_0 for each example in the third column of Table 1. Write the new representation of the data in terms of the *rules of thumb*, f_1 and f_2 , in the fourth and fifth columns of Table 1.
- (b) **[10 points]** Find the hypothesis given by the weak learner that minimizes the error ϵ for that distribution. Place this hypothesis as the heading to the sixth column of Table 1, and give its prediction for each example in that column.

i	Label	Hypothesis 1 (1st iteration)				Hypothesis 2 (2nd iteration)			
		D_0	$f_1 \equiv$ [$x > _$]	$f_2 \equiv$ [$y > _$]	$h_1 \equiv$ [$______$]	D_1	$f_1 \equiv$ [$x > _$]	$f_2 \equiv$ [$y > _$]	$h_2 \equiv$ [$______$]
(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	(10)
1	—								
2	—								
3	+								
4	—								
5	—								
6	—								
7	+								
8	—								
9	+								
10	+								

Table 1: Table for Boosting results

- (c) **[10 points]** Now compute D_1 for each example, find the new best weak learners f_1 and f_2 , and select hypothesis that minimizes error on this distribution, placing these values and predictions in the seventh to tenth columns of Table 1.
- (d) **[10 points]** Write down the final hypothesis produced by AdaBoost.

What to submit: Fill out Table 1 as explained, show computation of α and $D_1(i)$, and give the final hypothesis, H_{final} .

2 Multi-class classification - 60 points

Consider a multi-class classification problem with k class labels $\{1, 2, \dots, k\}$. Assume that we are given m examples, labeled with one of the k class labels. Assume, for simplicity, that we have m/k examples of each type.

Assume that you have a learning algorithm L that can be used to learn Boolean functions. (E.g., think about L as the Perceptron algorithm). We would like to explore several ways to develop learning algorithms for the multi-class classification problem.

There are two schemes to use the algorithm L on the given data set, and produce a multi-class classification:

- **One vs. All:** For every label $i \in [1, k]$, a classifier is learned over the following data set: the examples labeled with the label i are considered “positive”, and examples labeled with any other class $j \in [1, k], j \neq i$ are considered “negative”.
- **All vs. All:** For every pair of labels $\langle i, j \rangle$, a classifier is learned over the following data set: the examples labeled with one class $i \in [1, k]$ are considered “positive”, and those labeled with the other class $j \in [1, k], j \neq i$ are considered “negative”.

- (a) **[20 points]** For each of these two schemes, answer the following:
- How many classifiers do you learn?
 - How many examples do you use to learn each classifier within the scheme?
 - How will you decide the final class label (from $\{1, 2, \dots, k\}$) for each example?
 - What is the computational complexity of the training process?
- (b) **[5 points]** Based on your analysis above of two schemes individually, which scheme would you prefer? Justify.
- (c) **[5 points]** You could also use a KERNELPERCEPTRON for a two-class classification. We could also use the algorithm to learn a multi-class classification. Does using a KERNELPERCEPTRON change your analysis above? Specifically, what is the computational complexity of using a KERNELPERCEPTRON and which scheme would you prefer when using a KERNELPERCEPTRON?
- (d) **[10 points]** We are given a magical black-box binary classification algorithm (we don't know how it works, but it just does!) which has a learning time complexity of $O(dn^2)$, where n is the total number of training examples supplied (positive+negative) and d is the dimensionality of each example. What are the overall training time complexities of the all-vs-all and the one-vs-all paradigms, respectively, and which training paradigm is most efficient?
- (e) **[10 points]** We are now given another magical black-box binary classification algorithm (wow!) which has a learning time complexity of $O(d^2n)$, where n is the total number of training examples supplied (positive+negative) and d is the dimensionality of each example. What are the overall training time complexities of the all-vs-all and the one-vs-all paradigms, respectively, and which training paradigm is most efficient, when using this new classifier?
- (f) **[10 points]** Suppose we have learnt an all-vs-all multi-class classifier and now want to proceed to predicting labels on unseen examples.

We have learnt a simple linear classifier with a weight vector of dimensionality d for each of the $m(m-1)/2$ classes ($w_i^T x = 0$ is the simple linear classifier hyperplane for each $i = [1, \dots, m(m-1)/2]$)

We have two evaluation strategies to choose from. For each example, we can:

- **Counting:** Do all predictions then do a majority vote to decide class label
- **Knockout:** Compare two classes at a time, if one loses, never consider it again. Repeat till only one class remains.

What are the overall evaluation time complexities per example for Counting and Knockout, respectively?